

IMPORTING ALL THE LIBRARY

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
warnings.filterwarnings("ignore")

%matplotlib inline
```

```
In [2]: #loading the dataset from sklearn
from sklearn.datasets import load_boston
```

```
In [4]: boston=load_boston()
```

```
In [5]: boston
```

```

Out[5]: {'data': array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02,
4.9800e+00],
[2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
9.1400e+00],
[2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
4.0300e+00],
...,
[6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
5.6400e+00],
[1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
6.4800e+00],
[4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
7.8800e+00]]),
'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15. ,
18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9])),
'feature_names': array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RA
D',
'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7)'),
'DESCR': ".. _boston_dataset:\n\nBoston house prices dataset\n-----
--\n\n**Data Set Characteristics:** \n\n      :Number of Instances: 506 \n\n      :Number o

```

f Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.\n\n :Attribute Information (in order):\n - CRIM per capita crime rate by town\n - ZN proportion of residential land zoned for lots over 25,000 sq.ft.\n - INDUS proportion of non-retail business acres per town\n - CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)\n - NOX nitric oxides concentration (parts per 10 million)\n - RM average number of rooms per dwelling\n - AGE proportion of owner-occupied units built prior to 1940\n - DIS weighted distances to five Boston employment centres\n - RAD index of accessibility to radial highways\n - TAX full-value property-tax rate per \$10,000\n - PTRATIO pupil-teacher ratio by town\n - B $1000(B_k - 0.63)^2$ where B_k is the proportion of black people by town\n - LSTAT % lower status of the population\n - MEDV Median value of owner-occupied homes in \$1000's\n\n :Missing Attribute Values: None\n\n :Creator: Harrison, D. and Rubinfeld, D.L.\n\nThis is a copy of UCI ML housing dataset.\n<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>\n\n\nThis dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.\n\nThe Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.\n\nThe Boston house-price data has been used in many machine learning papers that address regression problems. \n\n.. topic:: References\n\n - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.\n - Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.\n",\n 'filename': 'boston_house_prices.csv',\n 'data_module': 'sklearn.datasets.data'}

In [7]: `print(boston)`

```

{'data': array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02,
4.9800e+00],
[2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
9.1400e+00],
[2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
4.0300e+00],
...,
[6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
5.6400e+00],
[1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
6.4800e+00],
[4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
7.8800e+00]]), 'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1,
16.5, 18.9, 15. ,
18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9])), 'feature_name
s': array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7'), 'DESCR': ".. _boston_dataset:\n\nB
oston house prices dataset\n-----\n\n**Data Set Characteristics:**
\n\n :Number of Instances: 506 \n\n :Number of Attributes: 13 numeric/categorical
predictive. Median Value (attribute 14) is usually the target.\n\n :Attribute Informa

```

tion (in order):\n - CRIM per capita crime rate by town\n - ZN proportion of residential land zoned for lots over 25,000 sq.ft.\n - INDUS proportion of non-retail business acres per town\n - CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)\n - NOX nitric oxides concentration (parts per 10 million)\n - RM average number of rooms per dwelling\n - AGE proportion of owner-occupied units built prior to 1940\n - DIS weighted distances to five Boston employment centres\n - RAD index of accessibility to radial highways\n - TAX full-value property-tax rate per \$10,000\n - PTRATIO pupil-teacher ratio by town\n - B 1000(Bk - 0.63)^2 where Bk is the proportion of black people by town\n - LSTAT % lower status of the population\n - MEDV Median value of owner-occupied homes in \$1000's\n\n:Missing Attribute Values: None\n\n:Creator: Harrison, D. and Rubinfeld, D.L.\n\nThis is a copy of UCI ML housing dataset.\n<https://archive.ics.uci.edu/ml/machine-learning-database/housing/>\n\nThis dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.\n\nThe Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.\n\nThe Boston house-price data has been used in many machine learning papers that address regression problems. \n\n.. topic:: References\n\n - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.\n - Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.\n", 'filename': 'boston_house_prices.csv', 'data_module': 'sklearn.datasets.data'}

In [6]: `boston.keys()`

Out[6]: `dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename', 'data_module'])`

In [9]: `print(boston.DESCR)`

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of black people by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
In [10]: print(boston.data)
```

```

[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]

```

```
In [11]: print(boston.target)
```

```

[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.  18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21.  12.7 14.5 13.2 13.1 13.5 18.9 20.  21.  24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20.  16.6 14.4 19.4 19.7 20.5 25.  23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16.  22.2 25.  33.  23.5 19.4 22.  17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20.  20.8 21.2 20.3 28.  23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22.  22.9 25.  20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22.  20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18.  14.3 19.2 19.6 23.  18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14.  14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17.  15.6 13.1 41.3 24.3 23.3 27.  50.  50.  50.  22.7 25.  50.  23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50.  32.  29.8 34.9 37.  30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.  22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25.  23.3 28.7 21.5 23.  26.7 21.7 27.5 30.1
 44.8 50.  37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.  24.  25.1 31.5
 23.7 23.3 22.  20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44.  50.  36.  30.1 33.8 43.1 48.8 31.  36.5 22.8
 30.7 50.  43.5 20.7 21.1 25.2 24.4 35.2 32.4 32.  33.2 33.1 29.1 35.1
 45.4 35.4 46.  50.  32.2 22.  20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29.  24.8 22.  26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21.  23.8 23.1
 20.4 18.5 25.  24.6 23.  22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19.  18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25.  19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50.  50.  50.  50.  50.  13.8 13.8 15.  13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3  8.8  7.2 10.5  7.4 10.2 11.5 15.1 23.2
  9.7 13.8 12.7 13.1 12.5  8.5  5.  6.3  5.6  7.2 12.1  8.3  8.5  5.
 11.9 27.9 17.2 27.5 15.  17.2 17.9 16.3  7.  7.2  7.5 10.4  8.8  8.4
 16.7 14.2 20.8 13.4 11.7  8.3 10.2 10.9 11.  9.5 14.5 14.1 16.1 14.3
 11.7 13.4  9.6  8.7  8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13.  13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20.  16.4 17.7
 19.5 20.2 21.4 19.9 19.  19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12.  14.6 21.4 23.  23.7 25.  21.8 20.6 21.2 19.1 20.6 15.2  7.
  8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22.  11.9]

```

```
In [12]: print(boston.feature_names)
```

```

['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']

```

```
In [13]: ##let's Prepare the data frame
```

```
In [146... #organsing the dataset
dataset=pd.DataFrame(boston.data,columns=boston.feature_names)
```

```
In [15]: dataset.head()
```

Out[15]:		CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
	0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
	1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
	2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
	3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
	4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [147]: #adding the dependent or target feature
dataset["Price"]=boston.target
```

```
In [17]: dataset.head()
```

Out[17]:		CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
	0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
	1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
	2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
	3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
	4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

```
In [19]: #basic eda
```

```
In [20]: dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS        506 non-null    float64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    float64
 9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  Price       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

```
In [21]: dataset.describe().T
```


Out[21]:		count	mean	std	min	25%	50%	75%	max
	CRIM	506.0	3.613524	8.601545	0.00632	0.082045	0.25651	3.677083	88.9762
	ZN	506.0	11.363636	23.322453	0.00000	0.000000	0.00000	12.500000	100.0000
	INDUS	506.0	11.136779	6.860353	0.46000	5.190000	9.69000	18.100000	27.7400
	CHAS	506.0	0.069170	0.253994	0.00000	0.000000	0.00000	0.000000	1.0000
	NOX	506.0	0.554695	0.115878	0.38500	0.449000	0.53800	0.624000	0.8710
	RM	506.0	6.284634	0.702617	3.56100	5.885500	6.20850	6.623500	8.7800
	AGE	506.0	68.574901	28.148861	2.90000	45.025000	77.50000	94.075000	100.0000
	DIS	506.0	3.795043	2.105710	1.12960	2.100175	3.20745	5.188425	12.1265
	RAD	506.0	9.549407	8.707259	1.00000	4.000000	5.00000	24.000000	24.0000
	TAX	506.0	408.237154	168.537116	187.00000	279.000000	330.00000	666.000000	711.0000
	PTRATIO	506.0	18.455534	2.164946	12.60000	17.400000	19.05000	20.200000	22.0000
	B	506.0	356.674032	91.294864	0.32000	375.377500	391.44000	396.225000	396.9000
	LSTAT	506.0	12.653063	7.141062	1.73000	6.950000	11.36000	16.955000	37.9700
	Price	506.0	22.532806	9.197104	5.00000	17.025000	21.20000	25.000000	50.0000

In [22]: *#check the missing value*

In [23]: `dataset.isnull().sum()`

Out[23]:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
Price	0

dtype: int64

In [25]: `dataset.corr().T`

Out[25]:

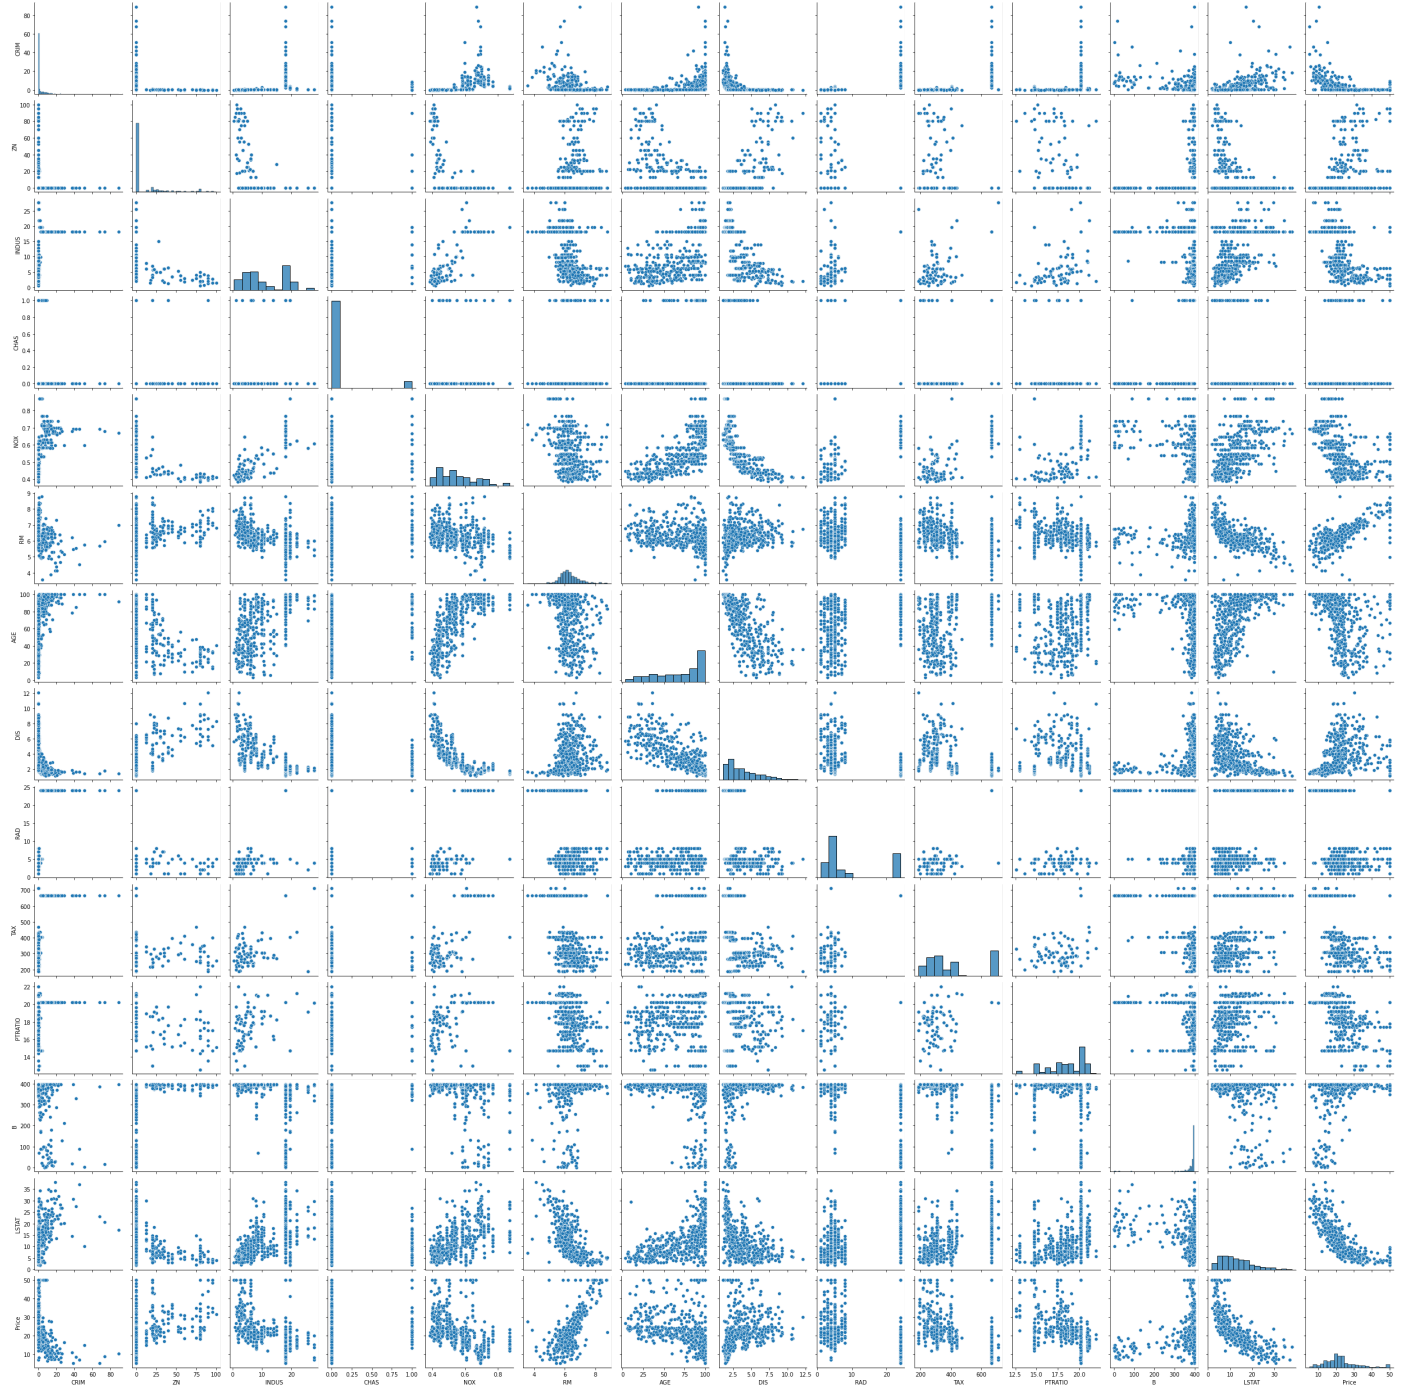
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.625505	0.5
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.311948	-0.3
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027	0.595129	0.7
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.007368	-0.0
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230	0.611441	0.6
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246	-0.209847	-0.2
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881	0.456022	0.5
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000	-0.494588	-0.5
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588	1.000000	0.9
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432	0.910228	1.0
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471	0.464741	0.4
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512	-0.444413	-0.4
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996	0.488676	0.5
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929	-0.381626	-0.4

In [26]:

sns.pairplot(dataset)


Out[26]:

<seaborn.axisgrid.PairGrid at 0x22b791347c0>




```
In [40]: sns.set(rc={"figure.figsize":(12,7)})
sns.heatmap(dataset.corr(),annot=True)
```

```
Out[40]: <AxesSubplot:>
```




```
In [35]: sns.set(rc={"figure.figsize":(10,8)})
plt.scatter(dataset["CRIM"],dataset["Price"])
plt.xlabel("CRIM")
plt.ylabel("Price")
```

```
Out[35]: Text(0, 0.5, 'Price')
```




```
In [42]: sns.set(rc={"figure.figsize":(10,8)})
plt.scatter(dataset["RM"],dataset["Price"])
plt.xlabel("RM")
plt.ylabel("Price")
```

```
Out[42]: Text(0, 0.5, 'Price')
```



```
In [44]: #for visulaization of best fit line
sns.regplot(x="RM",y="Price",data=dataset)
```


```
Out[44]: <AxesSubplot:xlabel='RM', ylabel='Price'>
```



In [45]: *#the shades region are ridge and laso regression to reduce the over fitting*


In [46]: *#for visulaization of best fit line*
`sns.regplot(x="LSTAT",y="Price",data=dataset)`

Out[46]: <AxesSubplot:xlabel='LSTAT', ylabel='Price'>




```
In [47]: #for visulaization of best fit line  
sns.regplot(x="CRIM",y="Price",data=dataset)
```

```
Out[47]: <AxesSubplot:xlabel='CRIM', ylabel='Price'>
```




```
In [48]: sns.boxplot(dataset["Price"])  
         #for dependent feature we dont do any things to outliers
```

```
Out[48]: <AxesSubplot:xlabel='Price'>
```

```
In [49]: sns.boxplot(dataset["CRIM"])
```

```
Out[49]: <AxesSubplot:xlabel='CRIM'>
```



In [50]: dataset.head()

Out[50]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

In [56]: *##Independent and Dependent Features*
x=dataset.iloc[:, :-1]
y=dataset.iloc[:, -1]

In [54]: y

Out[54]:

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2
	...
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: Price, Length: 506, dtype: float64

In [57]: x

Out[57]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

In [58]: *#after Eda we have to do train test split*

In [59]: from sklearn.model_selection import train_test_split

In [60]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=1

```
In [62]: X_train
```

```
Out[62]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
147	2.36862	0.0	19.58	0.0	0.871	4.926	95.7	1.4608	5.0	403.0	14.7	391.71	29.53
330	0.04544	0.0	3.24	0.0	0.460	6.144	32.2	5.8736	4.0	430.0	16.9	368.57	9.09
388	14.33370	0.0	18.10	0.0	0.700	4.880	100.0	1.5895	24.0	666.0	20.2	372.92	30.62
238	0.08244	30.0	4.93	0.0	0.428	6.481	18.5	6.1899	6.0	300.0	16.6	379.41	6.36
113	0.22212	0.0	10.01	0.0	0.547	6.092	95.4	2.5480	6.0	432.0	17.8	396.90	17.09
...
320	0.16760	0.0	7.38	0.0	0.493	6.426	52.3	4.5404	5.0	287.0	19.6	396.90	7.20
15	0.62739	0.0	8.14	0.0	0.538	5.834	56.5	4.4986	4.0	307.0	21.0	395.62	8.47
484	2.37857	0.0	18.10	0.0	0.583	5.871	41.9	3.7240	24.0	666.0	20.2	370.73	13.34
125	0.16902	0.0	25.65	0.0	0.581	5.986	88.4	1.9929	2.0	188.0	19.1	385.02	14.81
265	0.76162	20.0	3.97	0.0	0.647	5.560	62.8	1.9865	5.0	264.0	13.0	392.40	10.45

339 rows × 13 columns

```
In [64]: y_train
```

```
Out[64]:
```

147	14.6
330	19.8
388	10.2
238	23.7
113	18.7
...	...
320	23.8
15	19.9
484	20.6
125	21.4
265	22.8

Name: Price, Length: 339, dtype: float64

```
In [67]: X_train.shape
```

```
Out[67]: (339, 13)
```

```
In [68]: y_train.shape
```

```
Out[68]: (339,)
```

```
In [69]: X_test.shape
```

```
Out[69]: (167, 13)
```

```
In [70]: y_test.shape
```

```
Out[70]: (167,)
```

```
In [72]: #standardize the datasets
#feature scaling
```

```
In [73]: from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
In [76]: X_train=scaler.fit_transform(X_train)
```

```
In [77]: X_test=scaler.transform(X_test)  
#to avoid data leakage we use only transform not fit_transform
```

```
In [83]: X_train
```

```
Out[83]: array([[ -0.13641471,  -0.47928013,   1.16787606, ...,  -1.77731527,  
                0.39261401,   2.36597873],  
              [ -0.41777807,  -0.47928013,  -1.18043314, ...,  -0.75987458,  
                0.14721899,  -0.54115799],  
              [  1.31269177,  -0.47928013,   0.95517731, ...,   0.76628645,  
                0.19334986,   2.52100705],  
              ...,  
              [ -0.13520965,  -0.47928013,   0.95517731, ...,   0.76628645,  
                0.17012536,   0.06331026],  
              [ -0.40281114,  -0.47928013,   2.04022838, ...,   0.25756611,  
                0.32166792,   0.27238516],  
              [ -0.33104058,   0.34161649,  -1.07552092, ...,  -2.56351944,  
                0.39993132,  -0.34772815]])
```

```
In [78]: ##Model Training
```

```
In [79]: from sklearn.linear_model import LinearRegression
```

```
In [80]: regression=LinearRegression()
```

```
In [82]: regression
```

```
Out[82]: ▼ LinearRegression  
LinearRegression()
```

```
In [84]: #training the data  
regression.fit(X_train,y_train)
```

```
Out[84]: ▼ LinearRegression  
LinearRegression()
```

```
In [85]: #print the co-efficeints and the intercept  
print(regression.coef_)
```

```
[-1.29099218  1.60949999 -0.14031574  0.37201867 -1.76205329  2.22752218  
  0.32268871 -3.31184248  2.70288107 -2.09005699 -1.7609799   1.25191514  
 -3.83392028]
```

```
In [86]: print(regression.intercept_)
```

```
22.077286135693214
```

```
In [87]: #prediction for test data
```

```
In [88]: reg_pred=regression.predict(X_test)
```


```
In [89]: reg_pred
```

```
Out[89]: array([31.43849583, 31.98794389, 30.99895559, 22.31396689, 18.89492791,
16.21371128, 35.9881236 , 14.81264582, 25.04500847, 37.12806894,
21.49110158, 30.88757187, 28.05752881, 34.05600093, 33.75791114,
40.63880011, 24.24023412, 23.41351375, 25.54158122, 21.34135664,
32.71699711, 17.88341061, 25.49549436, 25.01006418, 32.54102925,
20.48979076, 19.48816948, 16.92733183, 38.38530857, 0.36265208,
32.42715816, 32.15306983, 26.10323665, 23.79611814, 20.67497128,
19.69393973, 3.50784614, 35.26259797, 27.04725425, 27.66164435,
34.35132103, 29.83057837, 18.40939436, 31.56953795, 17.91877807,
28.50042742, 19.49382421, 21.69553078, 38.0954563 , 16.44490081,
24.58507284, 19.67889486, 24.53954813, 34.30610423, 26.74699088,
34.87803562, 21.06219662, 19.87980936, 18.68725139, 24.71786624,
19.96344041, 23.56002479, 39.57630226, 42.81994338, 30.37060855,
17.03737245, 23.83719412, 3.2425022 , 31.5046382 , 28.63779884,
18.49288659, 27.14115768, 19.67125483, 25.34222917, 25.05430467,
10.29463949, 38.96369453, 8.26774249, 18.52214761, 30.34082002,
22.87681099, 20.96680268, 20.04604103, 28.73415756, 30.81726786,
28.23002473, 26.28588806, 31.59181918, 22.13093608, -6.48201197,
21.53000756, 19.90826887, 24.96686716, 23.44746617, 19.28521216,
18.75729874, 27.40013804, 22.17867402, 26.82972 , 23.39779064,
23.9260607 , 19.16632572, 21.09732823, 11.01452286, 13.7692535 ,
20.74596484, 23.54892211, 14.04445469, 28.88171403, 15.77611741,
15.25195598, 22.429474 , 26.60737213, 28.88742175, 24.29797261,
18.26839956, 16.26943281, 17.40100292, 15.53131616, 21.27868825,
33.78464602, 30.00899396, 21.16115702, 13.95560661, 16.18475215,
29.30998858, 13.1866784 , 22.08393725, 24.34499386, 31.86829501,
33.45923602, 5.90671516, 35.20153265, 24.17614831, 17.54200544,
24.25032915, 28.44671354, 34.50123773, 6.33164665, 1.93565618,
28.40727267, 12.56461105, 18.31045646, 19.71015745, 5.50105857,
14.51366874, 37.193992 , 25.81821367, 23.31632083, 26.43254504,
11.38255141, 20.46224115, 35.27645709, 20.57841598, 11.48799917,
16.23913171, 24.56511742, 10.53131603, 15.07115005, 25.98488217,
11.2136222 , 11.695686 , 19.40437966, 19.58768384, 32.43800883,
22.66170871, 25.68576052])
```

Assumption of Linear Regression

```
In [93]: plt.scatter(y_test,reg_pred)
plt.xlabel("Test truth Data")
plt.ylabel("Test pred Data")
```

```
Out[93]: Text(0, 0.5, 'Test pred Data')
```



#1 truth point and your predicted point if you try to plot into scatter plot this should be in linear manner (X incearsing y is incresing)

```
In [95]: #calculate the residuals
```


```
In [96]: residuals=y_test-reg_pred
```

```
In [97]: residuals
```

```
Out[97]: 305    -3.038496
193     -0.887944
65      -7.498956
349     4.286033
151     0.705072
...
442     -1.004380
451     -4.387684
188     -2.638009
76      -2.661709
314     -1.885761
Name: Price, Length: 167, dtype: float64
```

```
In [100... sns.displot(residuals,kind="kde")
```

```
Out[100]: <seaborn.axisgrid.FacetGrid at 0x22b0bc80e50>
```




#if we try to plot residuals in dis plot it should be coming into normal/guassian distribution

UNIFORM DISRIBUTION

In [101... *#scatter plot with predictions and residuals*

In [106... *#uniform distribution*
#Homoscedasticity
`plt.scatter(reg_pred,residuals)`

Out[106]: `<matplotlib.collections.PathCollection at 0x22b104f4b80>`



Performance Metrics

```
In [107... from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test, reg_pred))
print(mean_absolute_error(y_test, reg_pred))
print(np.sqrt(mean_squared_error(y_test, reg_pred)))
```

```
27.100991709962493
3.520658529879791
5.205861284164465
```

```
In [108... #R^2 and Adjusted R^2
```

```
In [111... from sklearn.metrics import r2_score
score=r2_score(y_test, reg_pred)
print(score)
```

```
0.7165219393967555
```

```
In [114... #adjusted R^2
1- (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

```
Out[114]: 0.6924355682343882
```

Ridge Regression

```
In [115... from sklearn.linear_model import Ridge
ridge=Ridge()
```

```
In [116... #model training
ridge.fit(X_train, y_train)
```

```
Out[116]: ▼ Ridge
Ridge()
```

```
In [119... #model prediction
ridge_pred=ridge.predict(X_test)
```

```
In [120... ridge_pred
```



```
Out[120]: array([[31.32951625, 31.98180665, 30.96523995, 22.45112285, 18.93171888,
16.21770197, 35.96932532, 14.8453389 , 25.00644473, 37.08826243,
21.49615236, 30.86395535, 27.9880323 , 33.98239498, 33.72731108,
40.61743429, 24.27292247, 23.33888547, 25.52862017, 21.42716828,
32.68689234, 17.88582539, 25.50293435, 25.01797349, 32.58757636,
20.48521647, 19.51598666, 16.94098815, 38.35803356, 0.33567931,
32.44411299, 32.10347472, 26.13567232, 23.81384315, 20.64388179,
19.71829821, 3.56174179, 35.17319673, 27.02020897, 27.65038259,
34.3408154 , 29.77237182, 18.39828682, 31.55283209, 17.92580288,
28.51408759, 19.49631857, 21.65517408, 38.03589465, 16.47721333,
24.56300743, 19.66060562, 24.490545 , 34.33513167, 26.7462751 ,
34.83714079, 21.08524522, 19.88396747, 18.65820105, 24.71538111,
20.00248822, 23.58585608, 39.60689645, 42.79543819, 30.3548884 ,
17.07425788, 23.84421168, 3.23169724, 31.42539336, 28.75103892,
18.49739555, 27.14667811, 19.64621723, 25.28950017, 25.07871104,
10.32212282, 38.94009655, 8.26854141, 18.50624966, 30.39028455,
22.88702308, 21.08817927, 20.09060901, 28.70289649, 30.81533585,
28.22566424, 26.28189093, 31.61850553, 22.15784726, -6.42142112,
21.55950809, 19.89786415, 24.96571959, 23.47361425, 19.25709566,
18.80383821, 27.37954116, 22.19229114, 26.78224659, 23.40784376,
23.92754566, 19.18858516, 21.09794643, 10.90877661, 13.8058827 ,
20.78603584, 23.49652544, 14.19685075, 28.86443391, 15.85586096,
15.26402087, 22.3935837 , 26.6360939 , 28.87654523, 24.25975975,
18.26463183, 16.26557102, 17.44937859, 15.58602415, 21.2407358 ,
33.72594686, 30.0710014 , 21.17366551, 14.04587364, 16.21847821,
29.26644762, 13.18724919, 22.07232566, 24.34918815, 31.88230457,
33.34230018, 5.95941842, 35.14730418, 24.25694454, 17.55532023,
24.27022839, 28.4213874 , 34.47544702, 6.3238347 , 2.03912756,
28.40127604, 12.59079125, 18.32110122, 19.75915926, 5.51559383,
14.42137586, 37.15183113, 25.8605775 , 23.29888263, 26.39528404,
11.42000684, 20.48891462, 35.29528497, 20.61619917, 11.45777136,
16.36445822, 24.57014519, 10.51041916, 15.13830095, 26.01152356,
11.22987126, 11.70179781, 19.39451509, 19.59207236, 32.42949 ,
22.67098418, 25.68376364])
```

Performance Matrix

```
In [121... #cost function mse,mae,rmse
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,ridge_pred))
print(mean_absolute_error(y_test,ridge_pred))
print(np.sqrt(mean_squared_error(y_test,ridge_pred)))

27.076490001440607
3.5161044263484236
5.203507471066089
```

```
In [122... #R square Metrics
```

```
In [134... from sklearn.metrics import r2_score
score=r2_score(y_test,ridge_pred)
print(score)

0.716778228793379
```

```
In [135... #adusted R^2
1- (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)

Out[135]: 0.6927136338542543
```

Lasso Regression

```
In [124... from sklearn.linear_model import Lasso
lasso=Lasso()
```

```
In [126... lasso
```

```
Out[126]: ▼ Lasso
          Lasso()
```

```
In [127... lasso.fit(X_train,y_train)
```

```
Out[127]: ▼ Lasso
          Lasso()
```

```
In [129... lasso_pred=lasso_pred=lasso.predict(X_test)
```

```
In [130... lasso_pred
```

```
Out[130]: array([25.64194382, 29.81425297, 27.94324255, 27.55256464, 20.99640298,
18.74520609, 34.28217994, 15.93009427, 20.70883387, 34.07542731,
19.90502439, 26.60490365, 24.07990755, 29.92866139, 29.22037693,
36.40160499, 26.2514407 , 19.88117334, 23.967085 , 22.50869347,
30.87428332, 18.78300957, 23.92041383, 25.68996484, 32.43275786,
21.59346217, 20.77097939, 19.17145706, 34.09829244, 2.83421427,
30.5699873 , 29.29565261, 26.85558827, 25.25346658, 19.26477827,
19.73302762, 7.84289608, 29.77239449, 25.40207471, 25.60513357,
32.3846261 , 26.89227407, 18.03007537, 29.36340326, 18.91119501,
27.26813644, 20.46203931, 21.03622196, 34.39115891, 18.05973586,
23.67935365, 18.6389767 , 22.52686697, 32.78082702, 26.03741902,
30.39354515, 20.51475327, 20.94796259, 17.76992156, 24.71515119,
21.39562999, 22.87363803, 36.66878913, 37.88636344, 28.19838095,
17.67593653, 24.95639783, 5.16197744, 27.44103022, 33.73592095,
19.49826488, 28.55803328, 19.79798303, 21.76346312, 24.39060267,
10.72390653, 35.69760879, 8.87515321, 19.7342389 , 30.44923304,
23.5731452 , 25.969039 , 21.53759864, 26.56759106, 29.08985079,
27.11618087, 27.52241607, 31.13762774, 25.19009251, 1.294237 ,
24.63082225, 20.35958514, 25.20780706, 25.3136586 , 19.78586427,
20.81458041, 26.18578331, 20.65389528, 23.63766687, 22.07598695,
24.48709112, 20.92346851, 21.99207783, 11.28215423, 15.85225507,
22.37860011, 20.11691913, 17.8652717 , 27.92655648, 20.66334411,
16.27505897, 21.25931319, 25.72062784, 25.92038934, 22.43890052,
18.21105925, 13.73609426, 19.70671479, 19.16823222, 20.81634433,
28.98351303, 31.19474714, 20.47647933, 14.48614909, 18.32476473,
25.90552665, 15.20476004, 22.61771889, 24.11123534, 30.06699424,
27.18076864, 10.97444888, 29.31613765, 29.11544566, 18.4851049 ,
25.5516343 , 26.58527747, 29.79753531, 7.90378337, 13.9914479 ,
27.53958439, 15.07581206, 19.1010104 , 20.2136539 , 9.56591918,
12.35215313, 32.65105454, 27.04505972, 22.60389757, 23.38431025,
18.14439345, 21.46916879, 33.35149732, 21.68370567, 11.686795 ,
22.46129098, 24.462044 , 10.63445232, 17.62333713, 25.46239518,
13.61074004, 15.25814185, 19.3752538 , 20.11939071, 29.36108994,
22.66082095, 25.30493792])
```

```
In [131... from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,lasso_pred))
```

```
print(mean_absolute_error(y_test, lasso_pred))
print(np.sqrt(mean_squared_error(y_test, lasso_pred)))
```

```
32.16822537607396
3.9064325476573205
5.67170392175702
```

```
In [132... from sklearn.metrics import r2_score
score=r2_score(y_test, lasso_pred)
print(score)
```

```
0.6635183597615237
```

```
In [133... #adusted R^2
1 - (1 - score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

```
Out[133]: 0.6349284164732871
```

ElasticNet Regression

```
In [136... from sklearn.linear_model import ElasticNet
```

```
In [137... elastic=ElasticNet()
```

```
In [138... elastic
```

```
Out[138]: ▼ ElasticNet
ElasticNet()
```

```
In [139... elastic.fit(X_train,y_train)
```

```
Out[139]: ▼ ElasticNet
ElasticNet()
```

```
In [140... elastic_pred=elastic.predict(X_test)
```

```
In [141... elastic_pred
```

```
Out[141]: array([26.0417533 , 29.72847396, 28.13249256, 27.33126697, 20.42880538,
 17.74088482, 31.34694254, 16.67485774, 22.66361605, 32.11606238,
 20.44062928, 27.05265082, 24.30388496, 29.10453835, 29.42032134,
 34.87404662, 25.31690008, 21.08018038, 24.04009667, 22.78241695,
 28.62957505, 18.35172223, 23.50225053, 24.94025282, 31.31440303,
 21.87551246, 22.30554751, 18.38033279, 33.5961939 ,  5.07350586,
 31.03524275, 28.19235387, 27.2862085 , 24.92462838, 19.28719449,
 20.2043877 ,  9.65913955, 29.64752478, 24.48773946, 25.34376165,
 30.68019641, 26.22751049, 18.01125345, 29.21052894, 20.61959202,
 27.27830384, 19.56149084, 19.72195809, 33.16071763, 19.16416141,
 23.05862027, 18.66118548, 22.77766754, 31.26962741, 25.0249516 ,
 29.94893114, 20.8407824 , 19.87498778, 18.27542547, 22.76517295,
 20.81723461, 22.76805785, 34.51940602, 36.11020157, 27.24493161,
 18.27047552, 24.17101249,  7.04406772, 26.85508847, 32.20329184,
 18.80351309, 27.46659498, 19.30788561, 20.8108208 , 23.90532445,
 14.22176442, 33.73545716, 10.78055539, 20.93445818, 29.78172162,
 23.99677889, 25.93581443, 22.61951728, 26.7428785 , 28.14408623,
 25.89892069, 26.67011775, 30.84900884, 24.58079972,  2.73998551,
 24.21010745, 20.83883219, 25.05619448, 24.60834531, 21.10174986,
 22.49049602, 26.06687733, 22.19194864, 23.68670917, 22.18041458,
 25.06905312, 20.26607354, 23.08760718, 10.49569995, 17.33370695,
 22.43253387, 21.15234101, 17.76115425, 26.58899485, 21.8834536 ,
 16.2603945 , 20.64765689, 25.09688902, 26.39218729, 22.53942481,
 18.04368235, 17.14638997, 20.67683019, 18.81134056, 19.82073711,
 26.78232308, 30.65570208, 20.62008364, 17.08768963, 19.18709774,
 25.65086934, 15.16746519, 21.83976175, 23.88029317, 29.16441724,
 27.25789963, 10.94687118, 29.43663882, 28.05096936, 18.27555631,
 24.83638128, 26.62936138, 30.15776119,  9.24328778, 10.68301626,
 27.04760052, 14.96584806, 18.73604079, 20.41851801,  9.92593546,
 14.91131492, 31.67918212, 27.61903571, 22.57544206, 24.72549427,
 15.35714158, 22.50035819, 31.86456612, 21.15334479, 12.67869784,
 22.65361671, 24.53650868, 11.9779066 , 18.36583206, 24.65573814,
 13.97917875, 14.8116782 , 18.92727223, 19.69541499, 28.7493839 ,
 22.61167089, 24.82104593])
```

```
In [142... from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,elastic_pred))
print(mean_absolute_error(y_test,elastic_pred))
print(np.sqrt(mean_squared_error(y_test,elastic_pred)))
```

```
35.341543853934674
4.035696708769101
5.944875427957652
```

```
In [143... from sklearn.metrics import r2_score
score=r2_score(y_test,elastic_pred)
print(score)
```

```
0.6303252509112042
```

```
In [144... #adusted R^2
1- (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

```
Out[144]: 0.5989149781128098
```

```
In [ ]:
```